

```
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib notebook
import seaborn as sns
import numpy as np
%matplotlib inline
# %load ../_standard_import.txt
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import make_classification

from sklearn.preprocessing import scale
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

from scipy.cluster import hierarchy

pd.set_option('display.notebook_repr_html', False)

%matplotlib inline
plt.style.use('seaborn-white')

<ipython-input-297-76b734dff1de>:24: MatplotlibDeprecationWarning: The seaborn styles shipped by Matplotlib are deprecated since 3.6, as they no longer correspond to the styles
plt.style.use('seaborn-white')
```

```
#Loading data
df = pd.read_csv('/content/churn.csv')
df.head()
```

None 



```
df.tail()
```

None 



▼ 2.2.1 Number of Observations:

```
df.shape
```

(10000, 14)

▼ 2.2.2 Field names

```
df.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
       'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
       'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

▼ 2.2.3 Type of fields

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column            Non-Null Count  Dtype  
 ---  --  
 0   RowNumber        10000 non-null   int64  
 1   CustomerId      10000 non-null   int64  
 2   Surname          10000 non-null   object  
 3   CreditScore      10000 non-null   int64  
 4   Geography         10000 non-null   object  
 5   Gender            10000 non-null   object  
 6   Age               10000 non-null   int64  
 7   Tenure           10000 non-null   int64  
 8   Balance           10000 non-null   float64 
 9   NumOfProducts     10000 non-null   int64  
 10  HasCrCard        10000 non-null   int64  
 11  IsActiveMember    10000 non-null   int64  
 12  EstimatedSalary   10000 non-null   float64 
 13  Exited           10000 non-null   int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
table = df.describe(include='all')
table
```

None  

▼ 3.1 Data Processing/Cleaning

```
df.isna().sum()
```

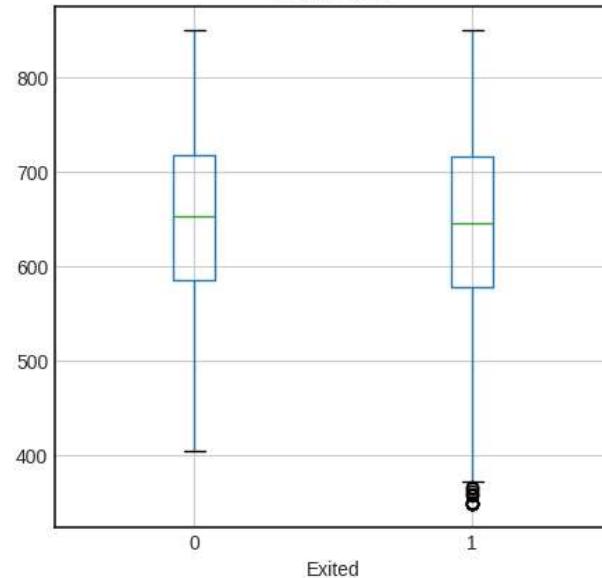
RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0

```
Gender      0  
Age        0  
Tenure     0  
Balance    0  
NumOfProducts 0  
HasCrCard   0  
IsActiveMember 0  
EstimatedSalary 0  
Exited      0  
dtype: int64
```

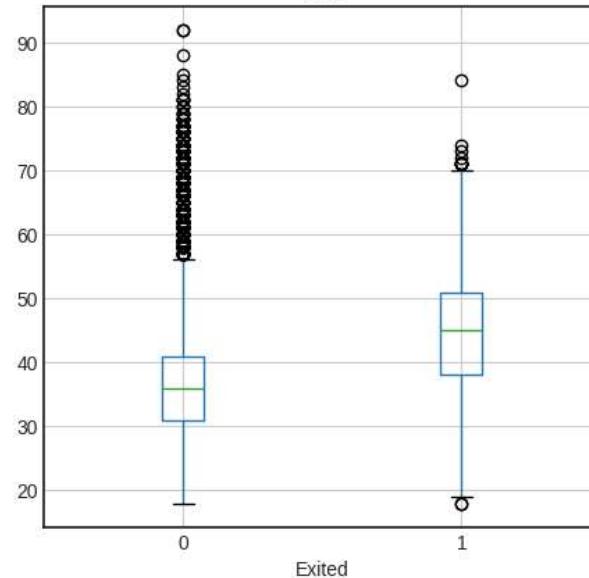
▼ 3.1.1 Outlier Detection

```
for i in df.columns.drop(['RowNumber', 'CustomerId', 'Surname', 'Geography', 'Gender', 'Exited']):  
    boxplot = df.boxplot(column=i, by = 'Exited', figsize=(5,5))
```

Boxplot grouped by Exited
CreditScore

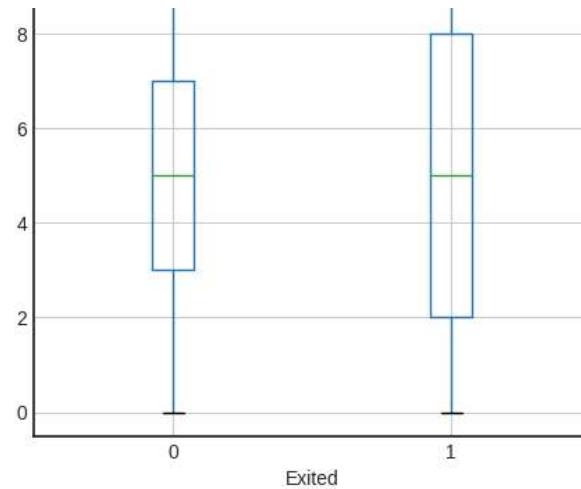


Boxplot grouped by Exited
Age

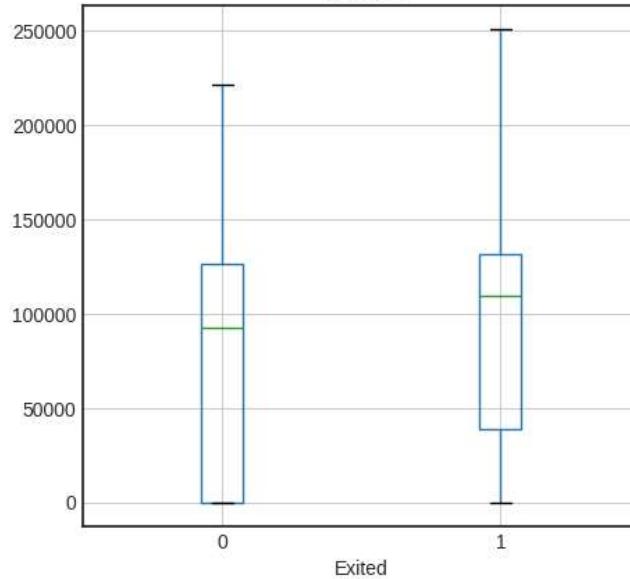


Boxplot grouped by Exited
Tenure

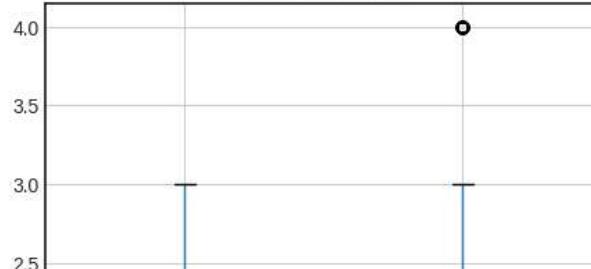


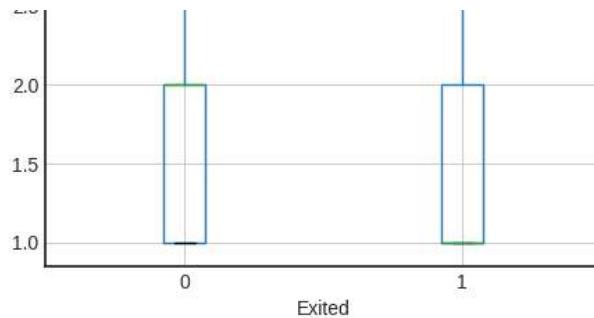


Boxplot grouped by Exited
Balance

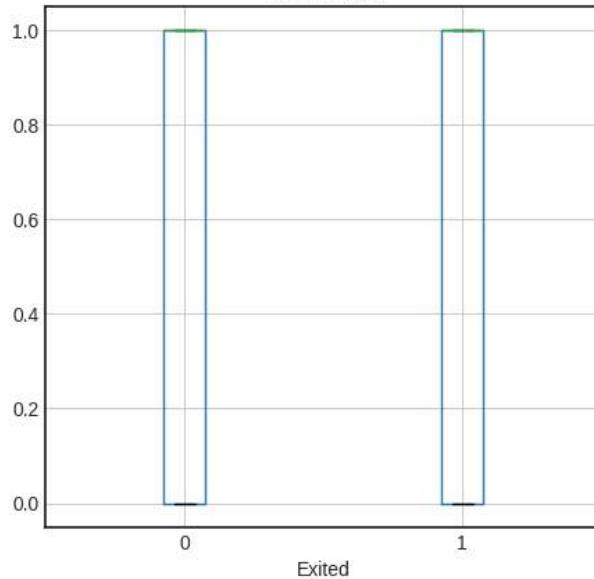


Boxplot grouped by Exited
NumOfProducts





Boxplot grouped by Exited
HasCrCard



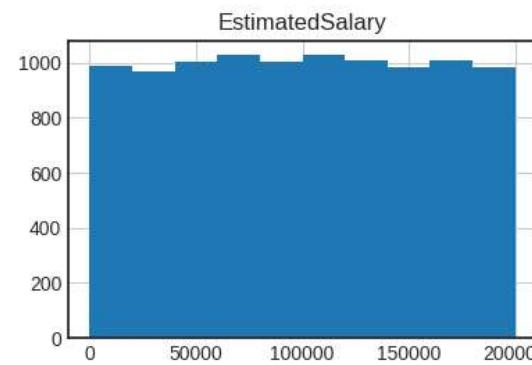
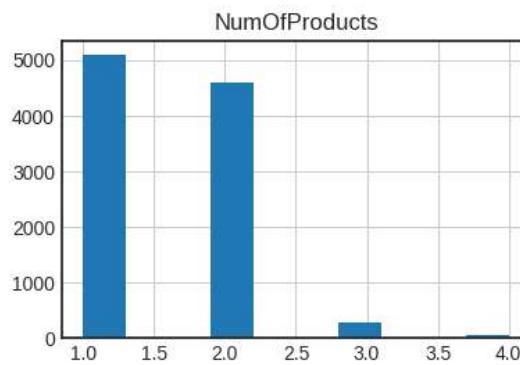
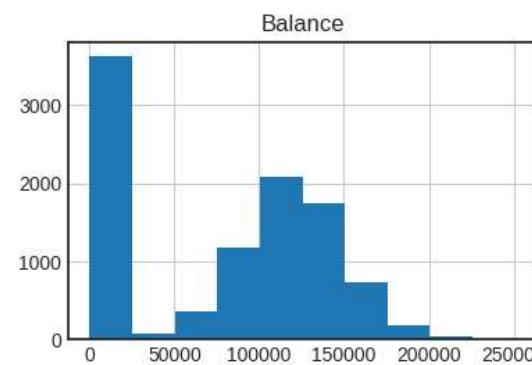
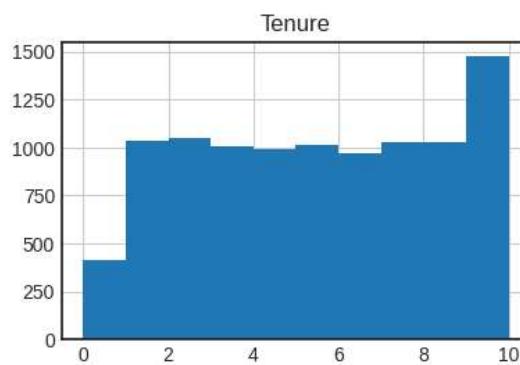
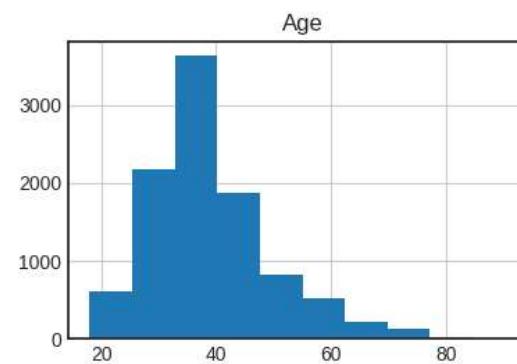
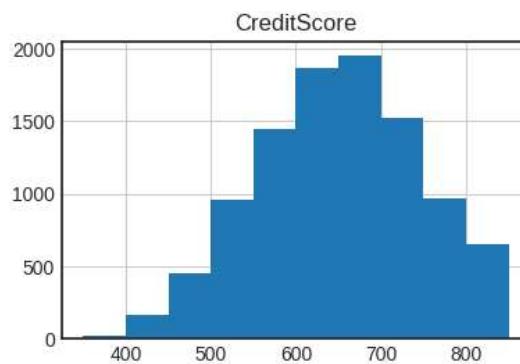
Boxplot grouped by Exited
IsActiveMember



▼ Histogram

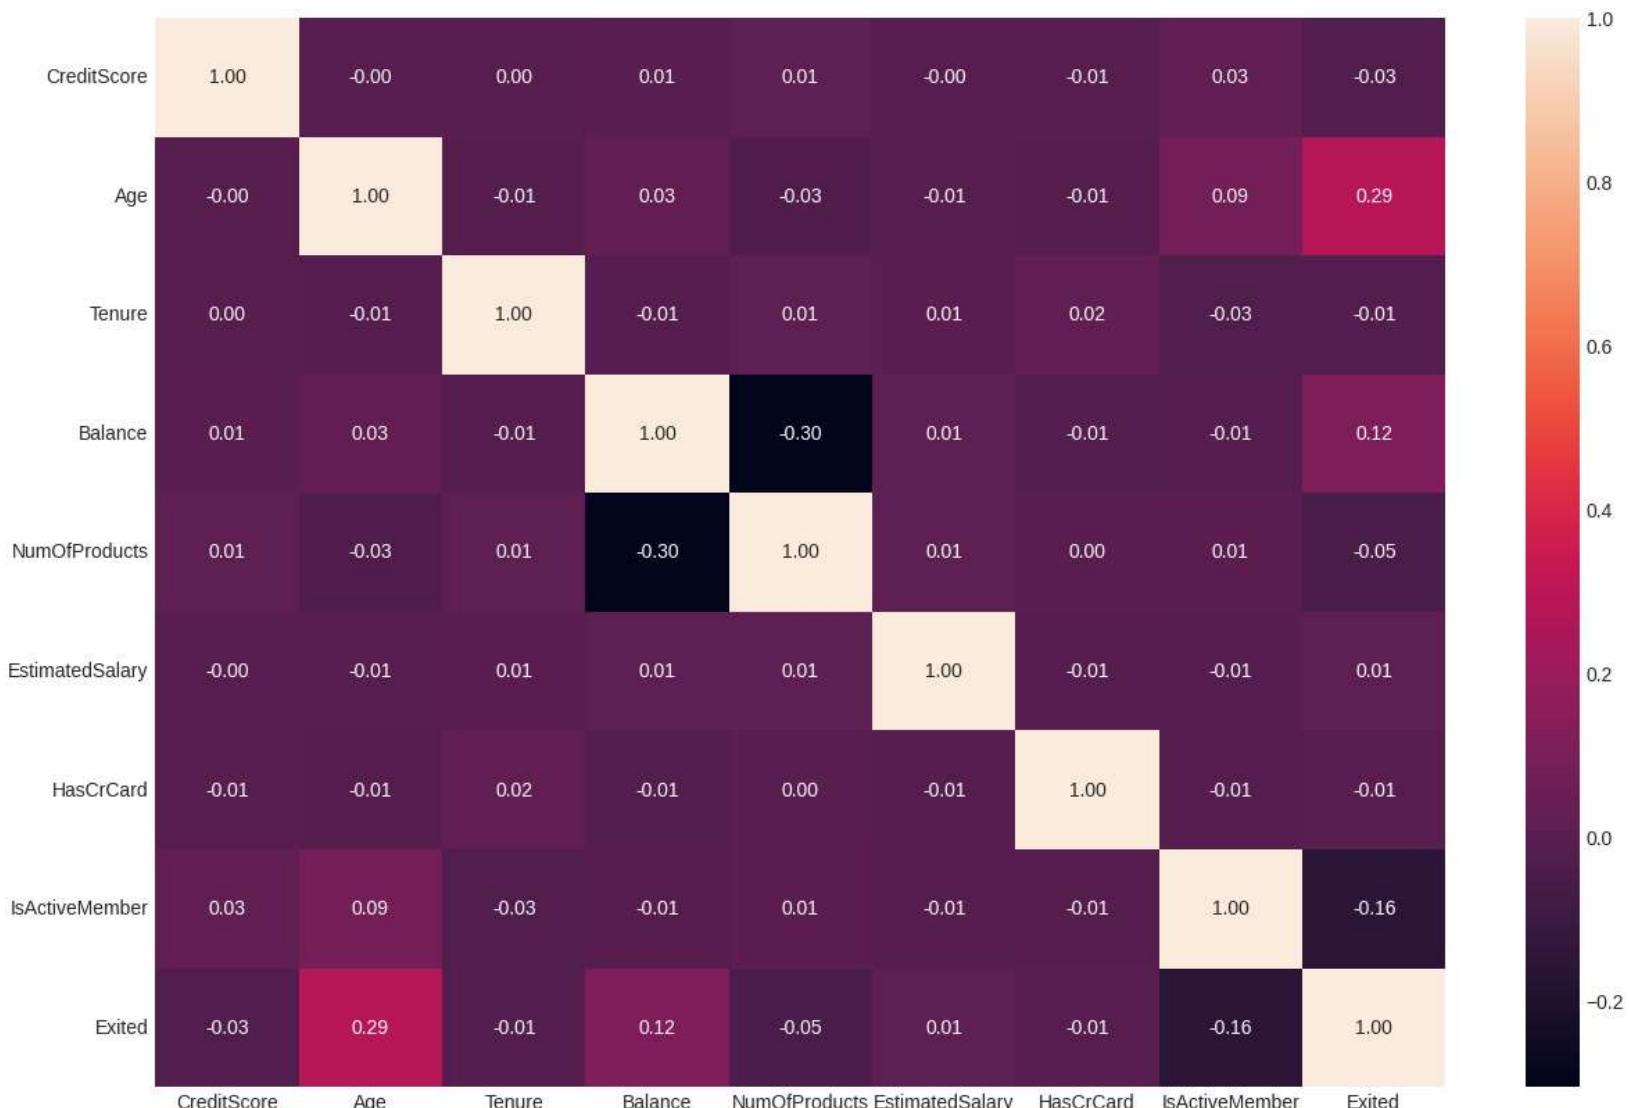


```
df.hist(column=['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary'], figsize=(10,10))
plt.show()
```

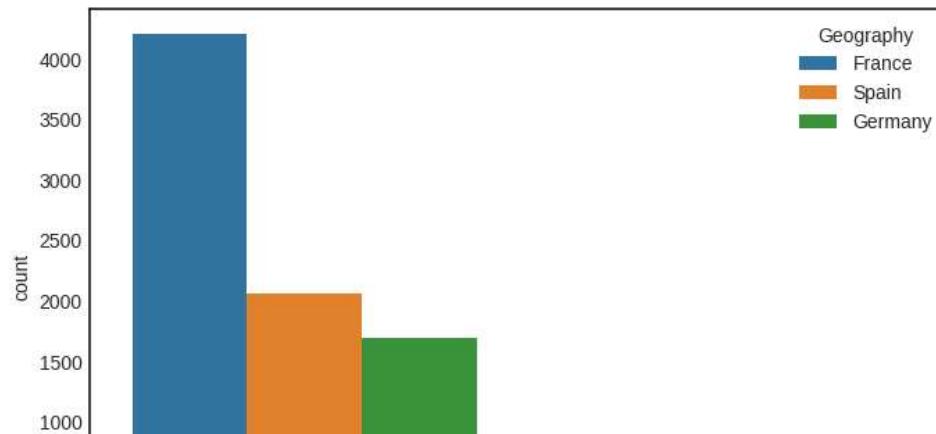


Correlation Matrix heatmap

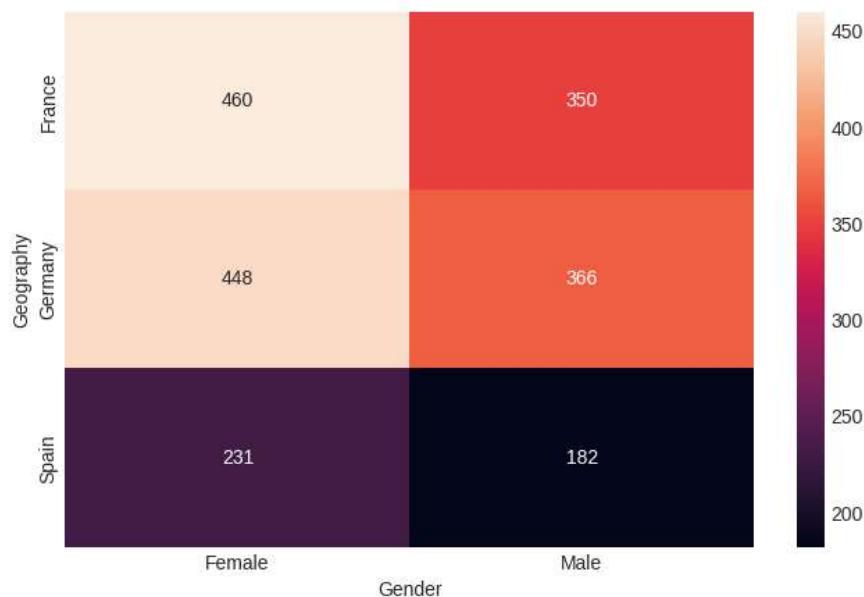
```
list1 = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary', 'HasCrCard', 'IsActiveMember', 'Exited']
plt.figure(figsize=(15,10))
sns.heatmap(df[list1].corr(), annot=True, fmt=".2f")
plt.show()
```



```
plt.figure(figsize=(8, 5))
sns.countplot(x='Exited', hue='Geography', data=df)
plt.show()
```



```
plt.figure(figsize=(8, 5))
crosstab = pd.crosstab(df['Geography'], df['Gender'], values=df['Exited'], aggfunc=np.sum)
sns.heatmap(crosstab, annot=True, fmt='d')
plt.show()
```



▼ Preprocessing & Feature Engineering

```
df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)
df.columns

Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
       'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
```

```
'Exited'],
dtype='object')

for i in range(len(df)):
    if df['EstimatedSalary'][i]<= 50000:
        df['EstimatedSalary'][i]= 'Poor'
    elif 50000 < df['EstimatedSalary'][i]<= 100000:
        df['EstimatedSalary'][i]= 'Good'
    elif 100000 < df['EstimatedSalary'][i]<= 200000:
        df['EstimatedSalary'][i]= 'Excellent'
    else:
        df['EstimatedSalary'][i]= "Top"

<ipython-input-311-f91b2aab9968>:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy.
df['EstimatedSalary'][i]= 'Excellent'

for i in range(len(df)):
    if df['Balance'][i]<= 25000:
        df['Balance'][i]= 'Poor'
    elif 25000 < df['Balance'][i]<= 100000:
        df['Balance'][i]= 'Good'
    elif 100000 < df['Balance'][i]<= 200000:
        df['Balance'][i]= 'Excellent'
    else:
        df['Balance'][i]= "Top"

<ipython-input-312-71ee8dca9bc5>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy.
df['Balance'][i]= 'Poor'

df['Gender'],df['Geography'],df['Balance'],df['EstimatedSalary']=df['Gender'].astype('category'),df['Geography'].astype('category'),df['Balance'].astype('category'),df['EstimatedSalary'].astype('category')
df['Gender']=df.Gender.cat.codes
df['Geography']=df.Geography.cat.codes
df['Balance']=df.Balance.cat.codes
df['EstimatedSalary']=df.EstimatedSalary.cat.codes

from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
df['CreditScore'] = scaler.fit_transform(df[['CreditScore']])
df['Age'] = scaler.fit_transform(df[['Age']])

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          -----          ----- 
 0   CreditScore  10000 non-null   float64
 1   Gender       10000 non-null   object  
 2   Geography    10000 non-null   object  
 3   Age          10000 non-null   float64
 4   Balance      10000 non-null   object  
 5   EstimatedSalary 10000 non-null   object  
 6   Education    10000 non-null   object  
 7   FamilySize   10000 non-null   int64  
 8   HasCrCard   10000 non-null   object  
 9   MaritalStatus 10000 non-null   object  
 10  Relationship 10000 non-null   object 
```

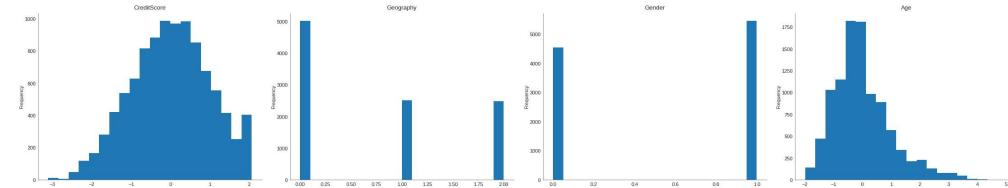
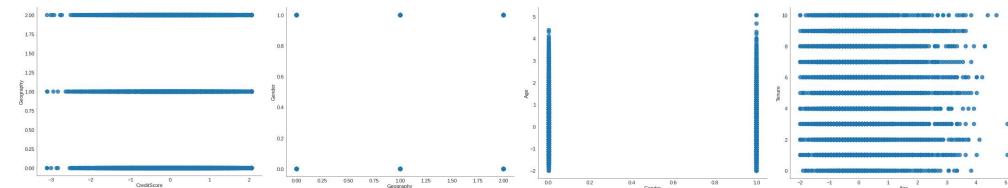
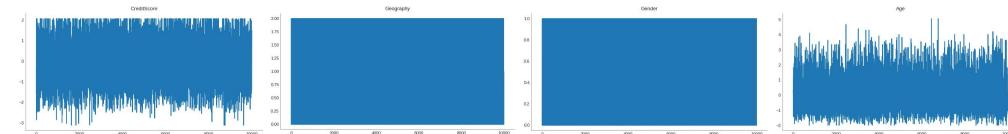
```

0 CreditScore      10000 non-null  float64
1 Geography        10000 non-null  int8
2 Gender           10000 non-null  int8
3 Age              10000 non-null  float64
4 Tenure           10000 non-null  int64
5 Balance          10000 non-null  int8
6 NumOfProducts    10000 non-null  int64
7 HasCrCard        10000 non-null  int64
8 IsActiveMember   10000 non-null  int64
9 EstimatedSalary  10000 non-null  int8
10 Exited          10000 non-null  int64
dtypes: float64(2), int64(5), int8(4)
memory usage: 586.1 KB

```

df

None

Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.**Distributions****2-d distributions****Values**

Error: Runtime no longer has a reference to this dataframe, please re-run this cell and try again.

Error: Runtime no longer has a reference to this dataframe, please re-run this cell and try again.

Error: Runtime no longer has a reference to this dataframe, please re-run this cell and try again.

```

predictors = ['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
              'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']
dependent_variable_name = ['Exited']

```

```

X = df[['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
         'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
X = scale(X)

```

```
X = pd.DataFrame(X)
X.head()
```

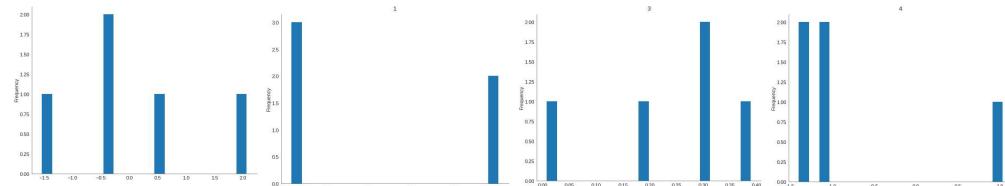
index	0	1	2	3	4	5
0	-0.32622142203674587	-0.9018862432746051	-1.0959875190286648	0.29351742289674687	-1.0417596792253108	1.2076506494805654
1	-0.4400359548576653	1.5150673766493	-1.0959875190286648	0.19816383219544517	-1.3875375865624426	0.11712270160214265
2	-1.5367941802228888	-0.9018862432746051	-1.0959875190286648	0.29351742289674687	1.03290776479748	-0.97340524627628
3	0.5015206348426682	-0.9018862432746051	-1.0959875190286648	0.007456650792841798	-1.3875375865624426	1.2076506494805654
4	2.063883767202562	1.5150673766493	-1.0959875190286648	0.38887101359804854	-1.0417596792253108	-0.97340524627628

Show 25 per page

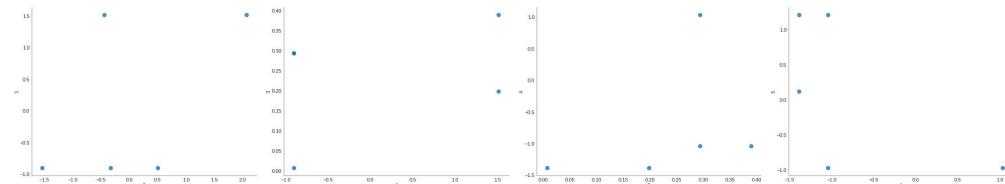


Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

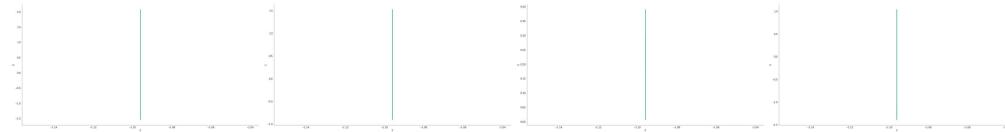
Distributions



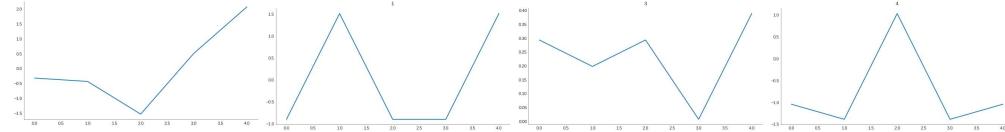
2-d distributions



Time series



Values



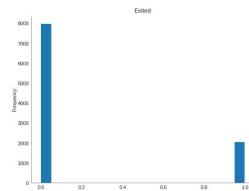
```
Y = df[['Exited']]
Y
```

1 to 25 of 10000 entries

index	Exited	
0	1	
1	0	
2	1	
3	0	
4	0	
5	1	
6	0	
7	1	
8	0	
9	0	
10	0	
11	0	
12	0	
13	0	
14	0	
15	0	
16	1	
17	0	
18	0	
19	0	
20	0	
21	0	
22	1	
23	0	
24	0	

Show per page Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

Distributions



Values



df['Exited'].value_counts()

```

0    7963
1    2037
Name: Exited, dtype: int64

```

```
from imblearn.under_sampling import RandomUnderSampler
```

```
rus = RandomUnderSampler(random_state=42)
X_train_rus, y_train_rus= rus.fit_resample(X, Y)

y_train_rus.value_counts()

Exited
0      2037
1      2037
dtype: int64
```

▼ Decision Tree

```
from sklearn.model_selection import train_test_split
np.random.seed(42)

X_train, X_test, y_train, y_test = train_test_split(X,Y, test_size=0.2, random_state=100)

print('The number of records in the training dataset is', X_train.shape[0])
print('The number of records in the test dataset is', X_test.shape[0])

The number of records in the training dataset is 8000
The number of records in the test dataset is 2000
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from collections import Counter

clf = DecisionTreeClassifier(criterion = 'entropy', random_state = 10, max_depth =8)
clf.fit(X_train, y_train)
```

```
▼          DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=8, random_state=10)
```

```
y_predic = clf.predict(X_test)
y_predic

array([0, 0, 0, ..., 0, 0, 0])

print("Accuracy of this model is:", accuracy_score(y_test, y_predic))

Accuracy of this model is: 0.8543333333333333

sco = accuracy_score(y_test, y_predic)*100
```

```
print("This model is", sco, "accurate, that is in acceptable " )  
  
This model is 85.43333333333332 accurate, that is in acceptable
```

```
print(classification_report(y_test, y_predic))
```

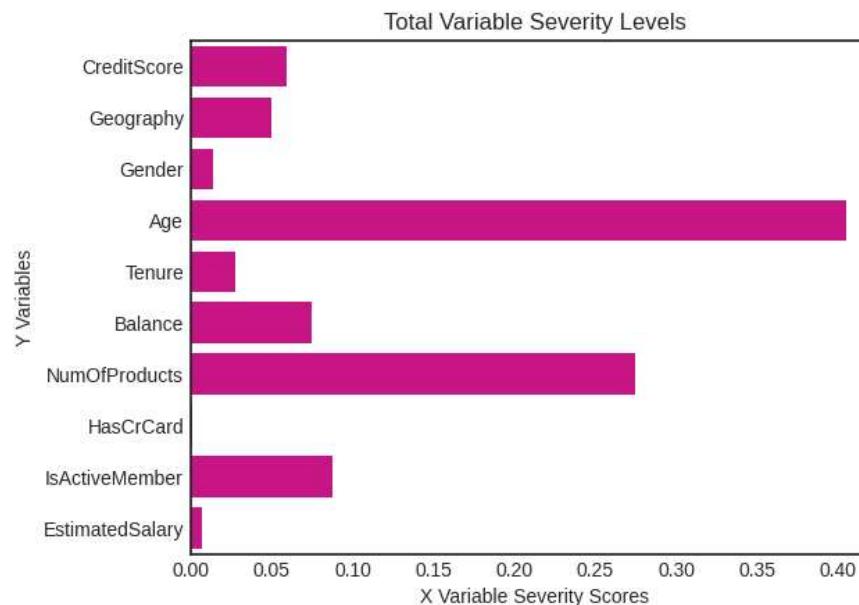
	precision	recall	f1-score	support
0	0.86	0.97	0.91	2416
1	0.76	0.37	0.49	584
accuracy			0.85	3000
macro avg	0.81	0.67	0.70	3000
weighted avg	0.84	0.85	0.83	3000

```
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier, export_graphviz  
import graphviz  
import pydot  
from IPython.display import Image, display  
  
feature_importances = clf.feature_importances_  
pd.DataFrame(data=feature_importances, index=predictors, columns = ['Importance']).sort_values(by=['Importance'], ascending=False)
```

index	Importance
Age	0.40521188612984704
NumOfProducts	0.27426100459028024

```
feature_importances_list = ['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']

sns.barplot(x = feature_importances, y = feature_importances_list, color='mediumvioletred', saturation=1)
plt.xlabel('X Variable Severity Scores')
plt.ylabel('Y Variables')
plt.title('Total Variable Severity Levels')
plt.show()
```



Random Forest

```
from sklearn.ensemble import RandomForestClassifier
RF=RandomForestClassifier(n_estimators=20)
RF.fit(X_train,y_train)

<ipython-input-371-55805171086b>:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example
  RF.fit(X_train,y_train)
  ▾      RandomForestClassifier
  RandomForestClassifier(n_estimators=20)
```

```
y_pred=RF.predict(X_test)
```

```
from sklearn import metrics
# Model Accuracy, how often is the classifier correct?
print("Accuracy of model:",metrics.accuracy_score(y_test, y_pred))

Accuracy of model: 0.8596666666666667
```

```
from sklearn.metrics import RocCurveDisplay

from sklearn.metrics import confusion_matrix, classification_report

confusion_matrix(y_test, y_pred)

array([[1535,    53],
       [ 235,  177]])
```

```
print(classification_report(y_test, y_pred))
```

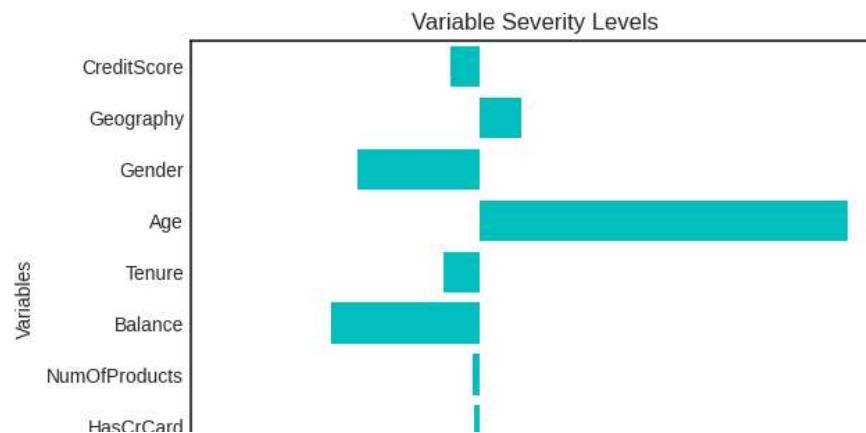
	precision	recall	f1-score	support
0	0.87	0.97	0.91	1588
1	0.77	0.43	0.55	412
accuracy			0.86	2000
macro avg	0.82	0.70	0.73	2000
weighted avg	0.85	0.86	0.84	2000

```
feature_importances = RF.feature_importances_
pd.DataFrame(data=feature_importances, index=predictors, columns = ['Importance']).sort_values(by=['Importance'], ascending=False)
```

None 


```
feature_importances_list = ['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']
```

```
sns.barplot(x = feature_importances, y = feature_importances_list, color='c', saturation=1)
plt.xlabel('Variable Severity Scores')
plt.ylabel('Variables')
plt.title('Variable Severity Levels')
plt.show()
```



▼ logistic Regression

```
from sklearn.linear_model import LogisticRegression
log_model = LogisticRegression(solver='liblinear', max_iter=50)
log_model.fit(X_train, y_train)

/usr/local/lib/python3.10/dist-packages/scikit-learn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the st
y = column_or_1d(y, warn=True)
  LogisticRegression
LogisticRegression(max_iter=50, solver='liblinear')
```

```
y_log_pred = log_model.predict(X_test)
y_log_pred.shape[0]
```

```
3000
```

```
print(y_log_pred)
```

```
[0 0 0 ... 0 0 0]
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
f1_score(y_test,y_log_pred)
```

```
0.27792207792207796
```

```
accuracy_score(y_test,y_log_pred)
```

```
0.8146666666666667
```

```
precision_score(y_test,y_log_pred)
```

```
0.5752688172043011
```

```
recall_score(y_test,y_log_pred)
```

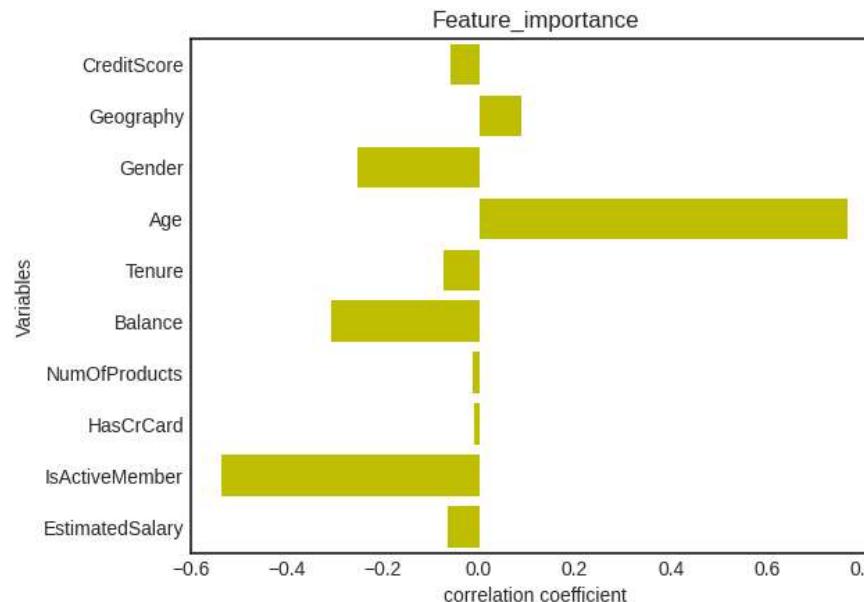
```
0.1832191780821918
```

```
print(classification_report(y_test,y_log_pred))
```

	precision	recall	f1-score	support
0	0.83	0.97	0.89	2416
1	0.58	0.18	0.28	584
accuracy			0.81	3000
macro avg	0.70	0.58	0.59	3000
weighted avg	0.78	0.81	0.77	3000

```
feature_importances = log_model.coef_[0]
pd.DataFrame(data=feature_importances, index=predictors,
              columns = ['Importance']).sort_values(by=['Importance'], ascending=False)
```

index	Importance
feature_importances_list = ['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']	
sns.barplot(x = feature_importances, y = feature_importances_list, color='y', saturation=10)	
plt.xlabel('correlation coefficient')	
plt.ylabel('Variables')	
plt.title('Feature_importance')	
plt.show()	



▼ KNN

```
from sklearn.neighbors import KNeighborsClassifier
KNN = KNeighborsClassifier(n_neighbors=8)

KNN.fit(X_train, y_train)

/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:215: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the code accordingly.
  return self._fit(X, y)
▼      KNeighborsClassifier
KNeighborsClassifier(n_neighbors=8)
```

```
y_pred=KNN.predict(X_test)

# Model Accuracy, how often is the classifier correct?
print("Accuracy of model:",metrics.accuracy_score(y_test, y_pred))

Accuracy of model: 0.838

confusion_matrix(y_test, y_pred)

array([[2345,    71],
       [ 415, 169]])
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.97	0.91	2416
1	0.70	0.29	0.41	584
accuracy			0.84	3000
macro avg	0.78	0.63	0.66	3000
weighted avg	0.82	0.84	0.81	3000

```
### Voting
```

```
from sklearn.ensemble import VotingClassifier
Classifier = VotingClassifier(estimators=[('DT', clf), ('RF', RF), ('LogReg', log_model), ('KNN', KNN)], voting='soft', flatten_transform=True)
eclf1 = Classifier.fit(X_train, y_train)
print(Classifier.score(X_test, y_test))

/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:99: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:134: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
0.8623333333333333
```

```
y_pred=Classifier.predict(X_test)
```

Hard and soft voting score

```
# importing libraries
from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import load_iris
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split

# train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,
                                                    Y,
                                                    test_size = 0.30,
                                                    random_state = 4)

# group / ensemble of models
estimator = []
estimator.append(('LR',
                  LogisticRegression(solver ='lbfgs',
                                      multi_class ='multinomial',
                                      max_iter = 200)))
estimator.append(('SVC', SVC(gamma ='auto', probability = True)))
estimator.append(('DTC', DecisionTreeClassifier()))

# Voting Classifier with hard voting
vot_hard = VotingClassifier(estimators = estimator, voting ='hard')
vot_hard.fit(X_train, y_train)
y_pred = vot_hard.predict(X_test)

# using accuracy_score metric to predict accuracy
score = accuracy_score(y_test, y_pred)
print("Hard Voting Score % d" % score)

# Voting Classifier with soft voting
vot_soft = VotingClassifier(estimators = estimator, voting ='soft')
vot_soft.fit(X_train, y_train)
y_pred = vot_soft.predict(X_test)

# using accuracy_score
score = accuracy_score(y_test, y_pred)
print("Soft Voting Score % d" % score)
```

```
↳ /usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:99: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
      y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:134: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
      y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
Hard Voting Score 0
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:99: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
      y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:134: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
      y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
Soft Voting Score 0
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.95	0.91	2416
1	0.68	0.46	0.55	584
accuracy			0.85	3000
macro avg	0.78	0.71	0.73	3000
weighted avg	0.84	0.85	0.84	3000

```
np.random.seed(10)
```

```
#set up plotting area
plt.figure(0).clf()
```

```
#fit logistic regression model and plot ROC curve
```

```
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Logistic Regression, AUC="+str(auc))
```

```
#fit KNN model and plot ROC curve
```

```
model = KNeighborsClassifier(n_neighbors=15)
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="KNN, AUC="+str(auc))
```

```
#fit Decision Tree model and plot ROC curve
```

```
model = DecisionTreeClassifier(criterion = 'entropy', random_state = 200, max_depth =8)
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Decision Tree, AUC="+str(auc))
```

```
#fit Random Forest model and plot ROC curve
```

```
model = RandomForestClassifier(n_estimators=400)
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Random Forest, AUC="+str(auc))
```

```
#fit Voting Classifier model and plot ROC curve
```

```
model = VotingClassifier(estimators=[('DT', clf), ('RF', RF), ('LogReg', log_model), ('KNN', KNN)],voting='soft',flatten_transform=True)
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
```

```
auc = round(metrics.roc_auc_score(y_test, y_preat), 4)
plt.plot(fpr,tpr,label="Voting Classifier, AUC="+str(auc))
```

```
#add legend
plt.legend()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the s
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:215: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please chang
    return self._fit(X, y)
<ipython-input-406-cc7550c66c2b>:32: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example
    model.fit(X_train, y_train)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:99: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:134: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
  y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
<matplotlib.legend.Legend at 0x7cf60da34f10>
```

